


#	List Size	#	List Size (k)	#	Threads	#	Threads (q)	#	Error	#	Time	#	qsort_time	#	Speedup	#	Efficiency	Speedup vs qsort
	268435456		28		1		0		0		38.7103		38.6755		1		1	0.9991010145
	268435456		28		2		1		0		19.9865		38.6143		1.936822355		0.9684111775	1.932019113
	268435456		28		4		2		0		10.6951		38.7942		3.619442548		0.9048606371	3.627287262
	268435456		28		16		4		0		2.9381		38.7122		13.17528335		0.8234552091	13.17593002
	268435456		28		64		6		0		1.4782		38.7987		26.18745772		0.4091790269	26.24726018
	268435456		28		256		8		0		1.7055		39.0104		22.69733216		0.08866145375	22.87329229
	268435456		28		1024		10		0		1.7479		39.0302		22.14674753		0.02162768313	22.32976715
	1048576		20		1		0		0		0.1084		0.1077		1		1	0.9935424354
	1048576		20		2		1		0		0.0583		0.1086		1.859348199		0.9296740995	1.862778731
	1048576		20		4		2		0		0.0319		0.1079		3.398119122		0.8495297806	3.382445141
	1048576		20		16		4		0		0.0115		0.1099		9.426086957		0.5891304348	9.556521739
	1048576		20		64		6		0		0.0096		0.1095		11.29166667		0.1764322917	11.40625
	1048576		20		256		8		0		0.0245		0.1097		4.424489796		0.01728316327	4.47755102
	1048576		20		1024		10		0		0.0679		0.1089		1.59646539		0.001559048233	1.603829161
	4096		12		1		0		0		0.0009		0.0003		1		1	0.3333333333
	4096		12		2		1		0		0.0008		0.0003		1.125		0.5625	0.375
	4096		12		4		2		0		0.0014		0.0003		0.6428571429		0.1607142857	0.2142857143
	4096		12		16		4		0		0.0016		0.0003		0.5625		0.03515625	0.1875
	4096		12		64		6		0		0.003		0.0003		0.3		0.0046875	0.1
	4096		12		256		8		0		0.0102		0.0003		0.08823529412		0.0003446691176	0.02941176471
	4096		12		1024		10		0		0.0496		0.0003		0.01814516129		0.00001771988407	0.006048387097

Speedup with respect to threads for k = 28


0	1
1	1.936822355
2	3.619442548
4	13.17528335
6	26.18745772
8	22.69733216
10	22.14674753

Efficiency	
0	1
1	0.9684111775
2	0.9048606371
4	0.8234552091
6	0.4091790269
8	0.08866145375
10	0.02162768313

vs. Speedup with respect to threads for k = 28



Efficiency




Here, the algorithm scales well, with peak speedup reaching 26x with 64 threads. After this point, there is diminishing returns, which is expected. Efficiency degrades as thread count increases, which is also expected.

Speedup with respect to threads for k = 20


0	1
1	1.859348199
2	3.398119122
4	9.426086957
6	11.29166667
8	4.424489796
10	1.59646539

Efficiency	
0	1
1	0.9296740995
2	0.8495297806
4	0.5891304348
6	0.1764322917
8	0.01728316327
10	0.001559048233

vs. Speedup with respect to threads for k = 20



Efficiency




Here, the algorithm doesn't scale quite as well as the previous problem, because there is less work to do. The max speed up still reaches a respectable 11x at 64 threads. Efficiency degrades similarly.

Speedup with respect to threads for k = 12


0	1
1	1.125
2	0.6428571429
4	0.5625
6	0.3
8	0.08823529412
10	0.01814516129

Efficiency	
0	1
1	0.5625
2	0.1607142857
4	0.03515625
6	0.0046875
8	0.0003446691176
10	0.00001771988407

vs. Speedup with respect to threads for k = 12



Efficiency



Here, performances is bad for anything more than two threads. Evidently when the list is so small, the sorting is trivial and multithreading overhead far exceeds single thread sort time. Efficiency expectedly drops off a cliff.

#	List Size	#	List Size (k)	#	Threads	#	Threads (q)	#	Error	#	Time	#	qsort_time	#	Speedup	#	Efficiency	Speedup vs qsort
part 3.																		
q = 6 seems to be the cut off point for diminishing returns, asm demonstrated in both k = 20 and k = 28.																		
In terms of varying q, it is up to the application, going beyond q seems to still yield good speed up in large lists, and in much larger lists the scaling will likely be much better, with larger q values being the cut off rather than q = 6																		
Based on the plots above and my accompanying analysis, the implementation scales well with large lists, while performing badly against small lists, which is all very expected behavior.																		